


Kubernetes Networks

How kubernetes networks work

David Chang

Linkernetworks
BackEnd, DevOps,
Docker, Kubernetes

dchang@linkernetworks.com



Outline

- Docker containers networks
- Containers communication in a Pod
- Pods cross different nodes
- Pod to Service

Docker Container Networks

- Bridge networks
communicate namespaces through bridge
- Host networks
use host's port, ip...
- Container networks

Docker Bridge Networks

```
4: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN  
    link/ether 02:42:89:14:7d:a5 brd ff:ff:ff:ff:ff:ff  
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0  
        valid_lft forever preferred_lft forever  
vagrant@master-1:~$
```

Docker daemon added a bridge, docker0 on host

Docker Bridge Networks

```
vagrant@master-1:~$ sudo docker attach busybox1
/ # ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
5: eth0@if6: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue
    link/ether 02:42:ac:11:00:02 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.2/16 brd 172.17.255.255 scope global eth0
        valid_lft forever preferred_lft forever
```

\$ docker run -itd --name busybox1 busy box

A eth0 inside container

Docker Bridge Networks

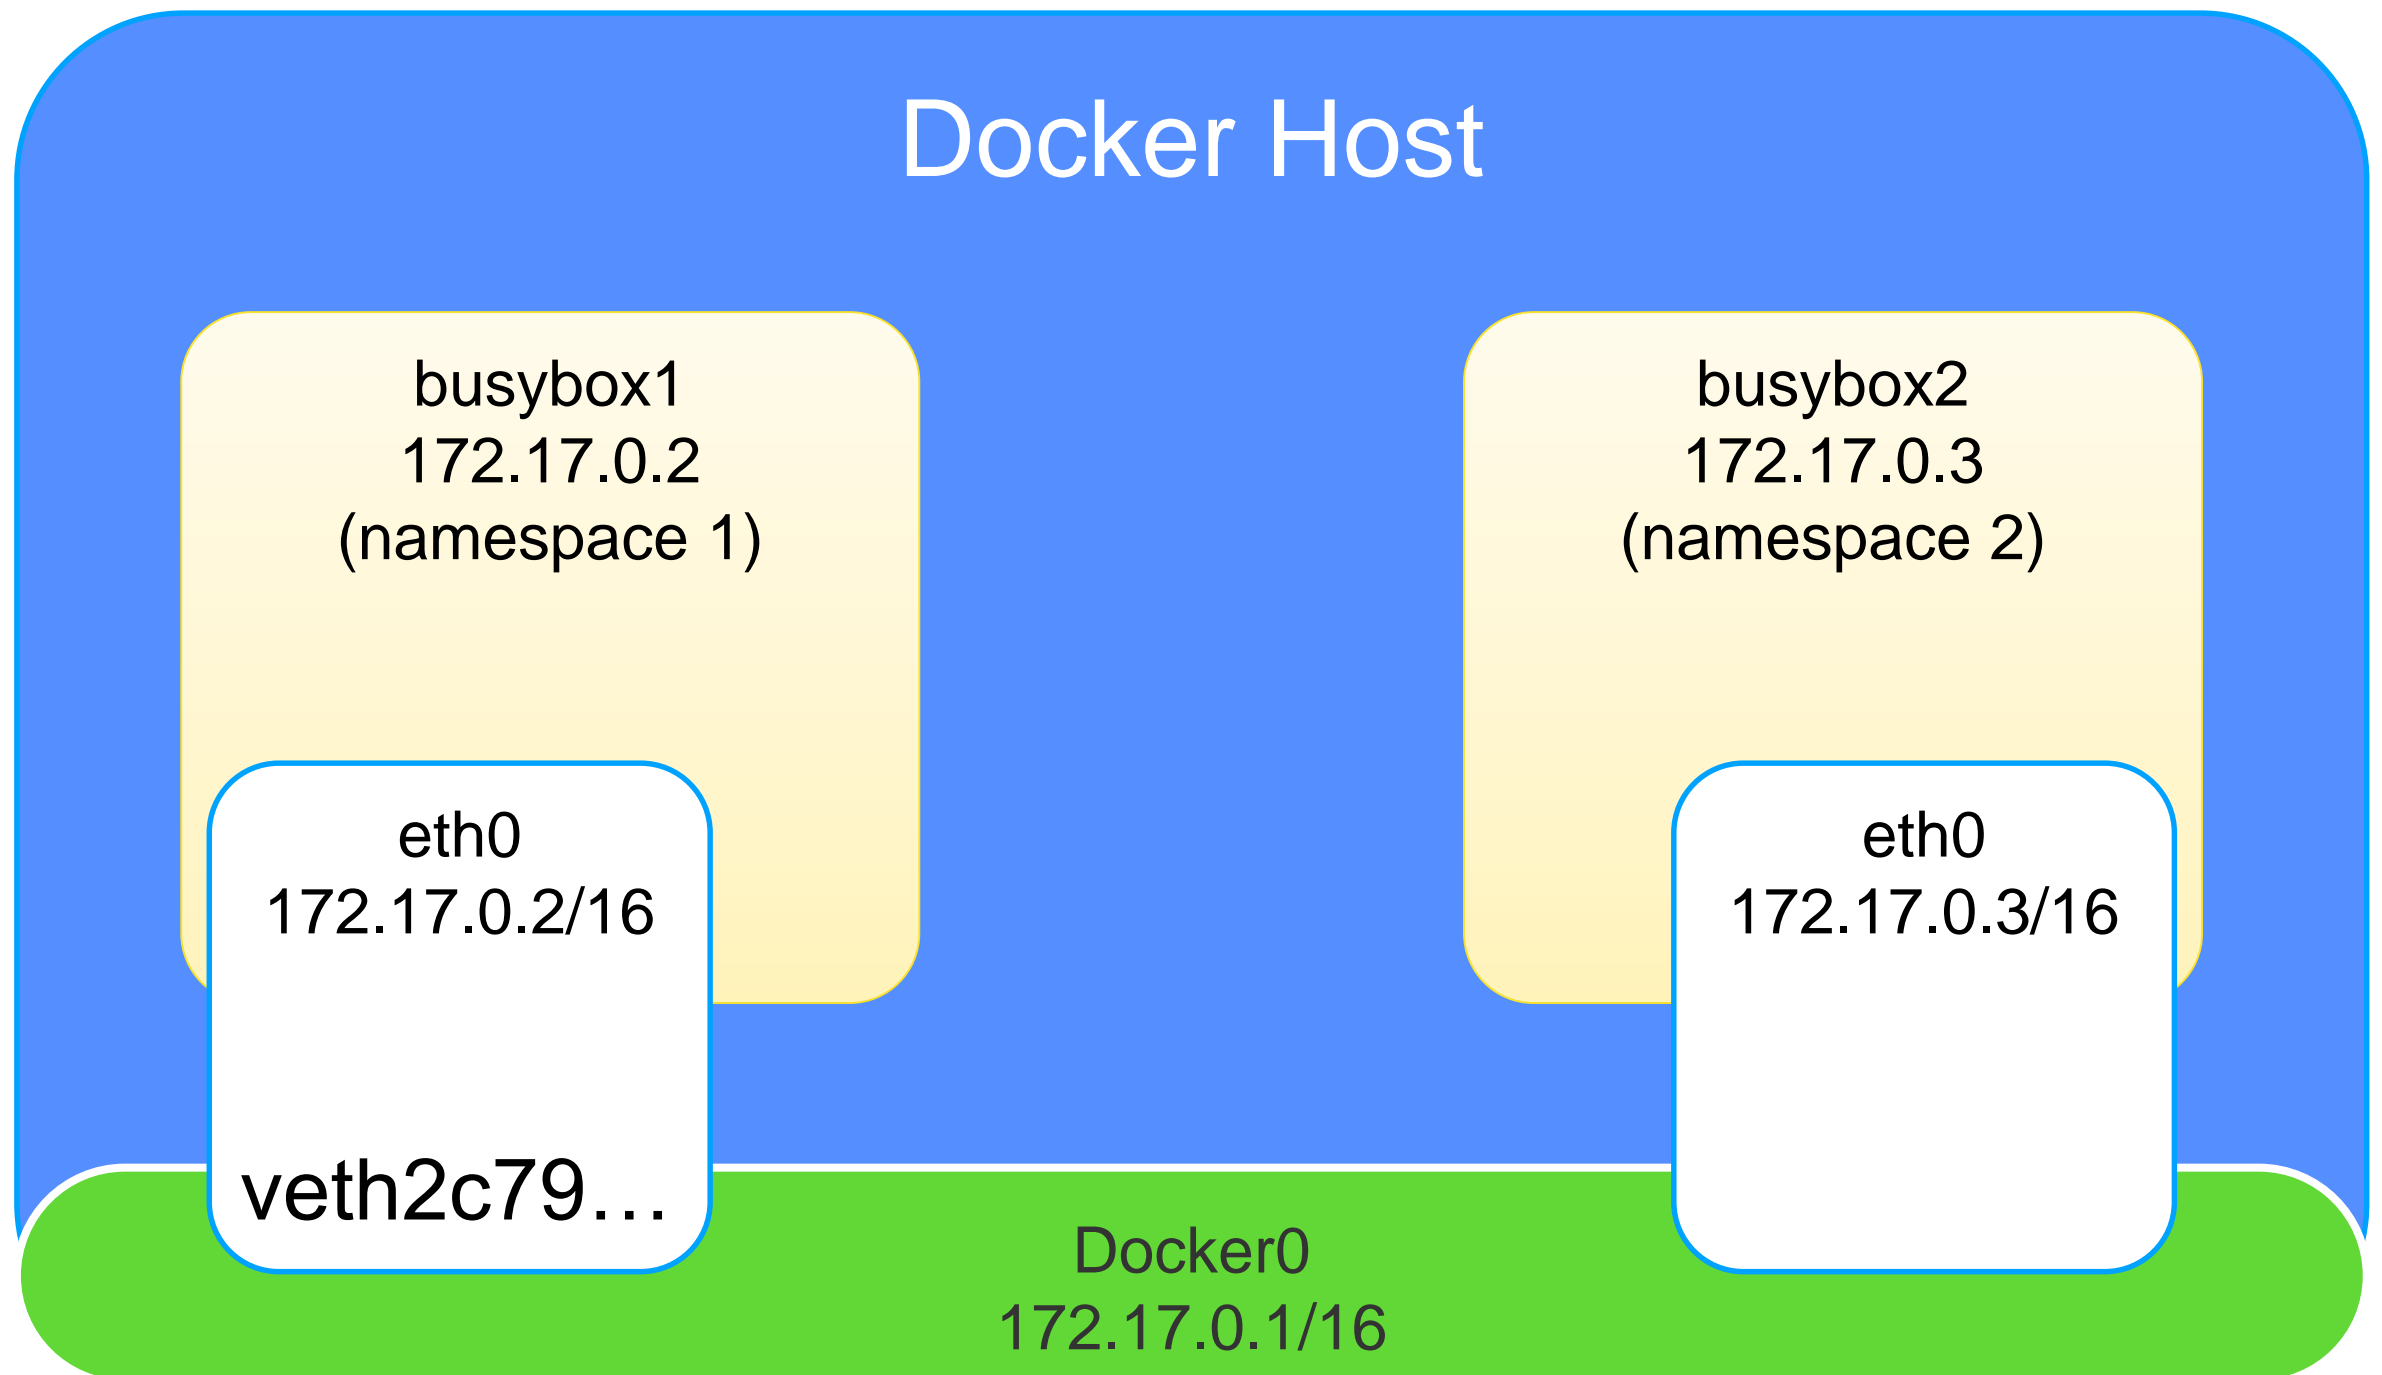
```
4: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue
    link/ether 02:42:89:14:7d:a5 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:89ff:fe14:7da5/64 scope link
        valid_lft forever preferred_lft forever
8: vethc2c792c@if7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
    link/ether 9e:d8:c0:c8:74:81 brd ff:ff:ff:ff:ff:ff link-netnsid
    inet6 fe80::9cd8:c0ff:fec8:7481/64 scope link
        valid_lft forever preferred_lft forever
10: veth8b23b0b@if9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc
    link/ether c2:5c:79:01:c1:68 brd ff:ff:ff:ff:ff:ff link-netnsid
    inet6 fe80::c05c:79ff:fe01:c168/64 scope link
        valid_lft forever preferred_lft forever
```

eth0@busybox1 - veth2c792c@if7 - docker0

Assign an ip from docker0 to eth@busybox1

<https://docs.docker.com/v17.09/engine/userguide/networking/#default-networks>

Docker Bridge Networks



Container has unique ip on a single docker host

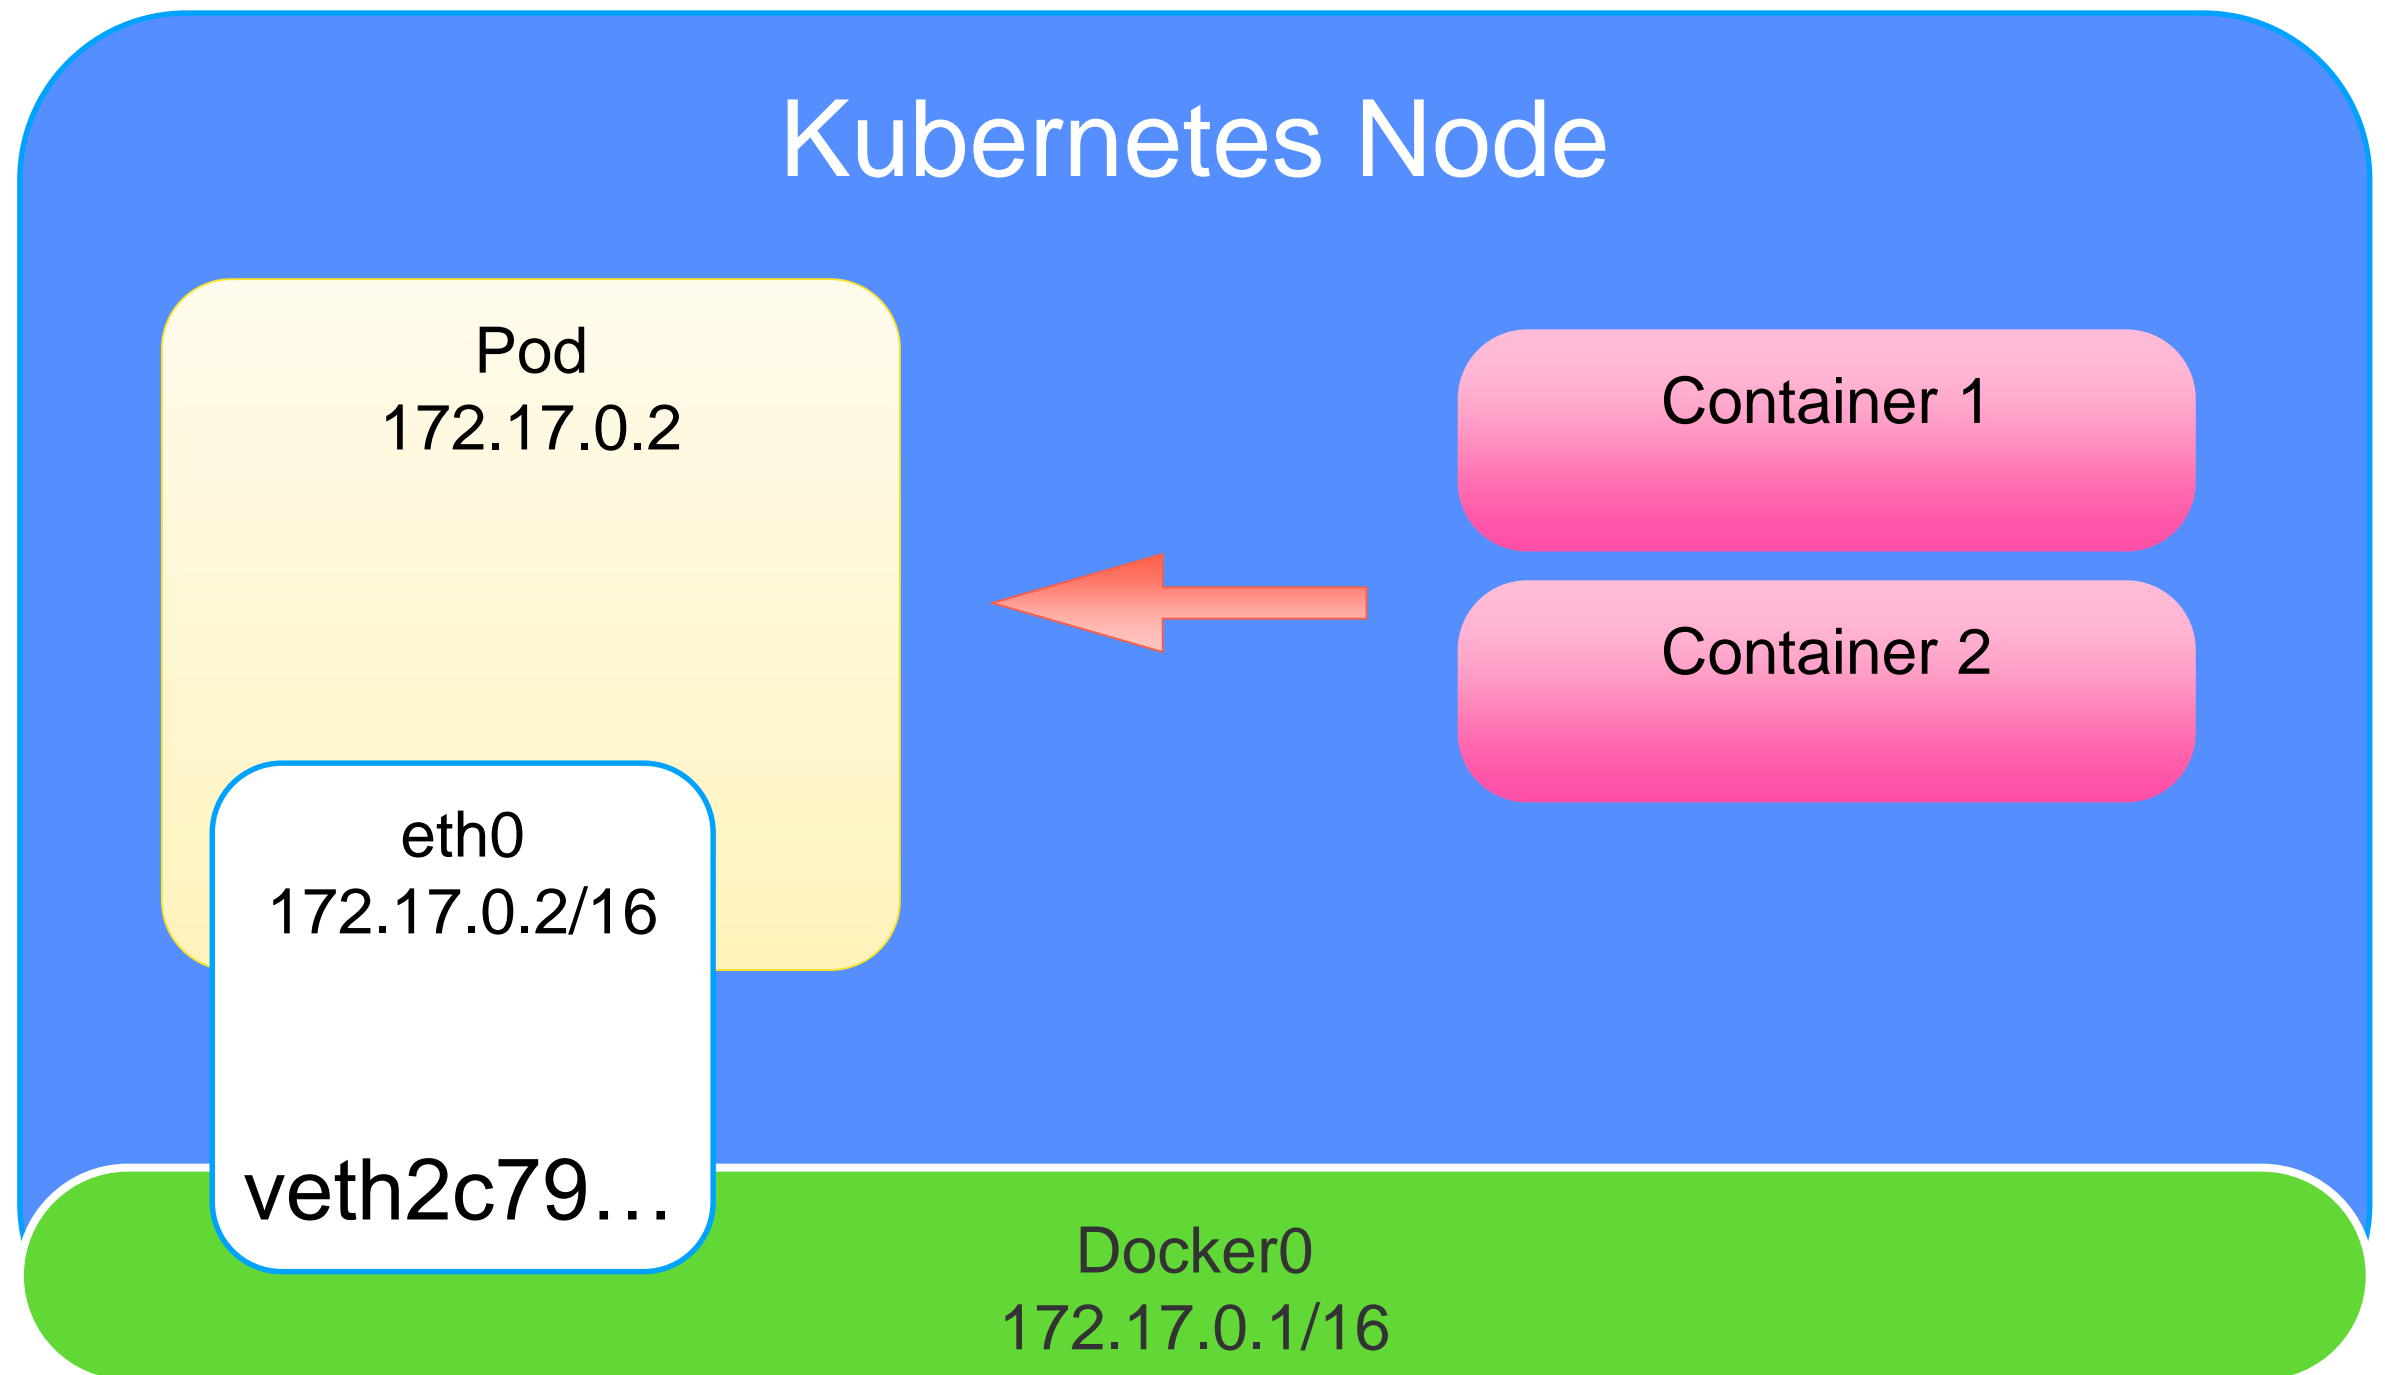
Kubernetes Model

- Docker container networks
- Containers communication in a Pod
- Pods cross different nodes
- Pod to Service

Kubernetes Networks

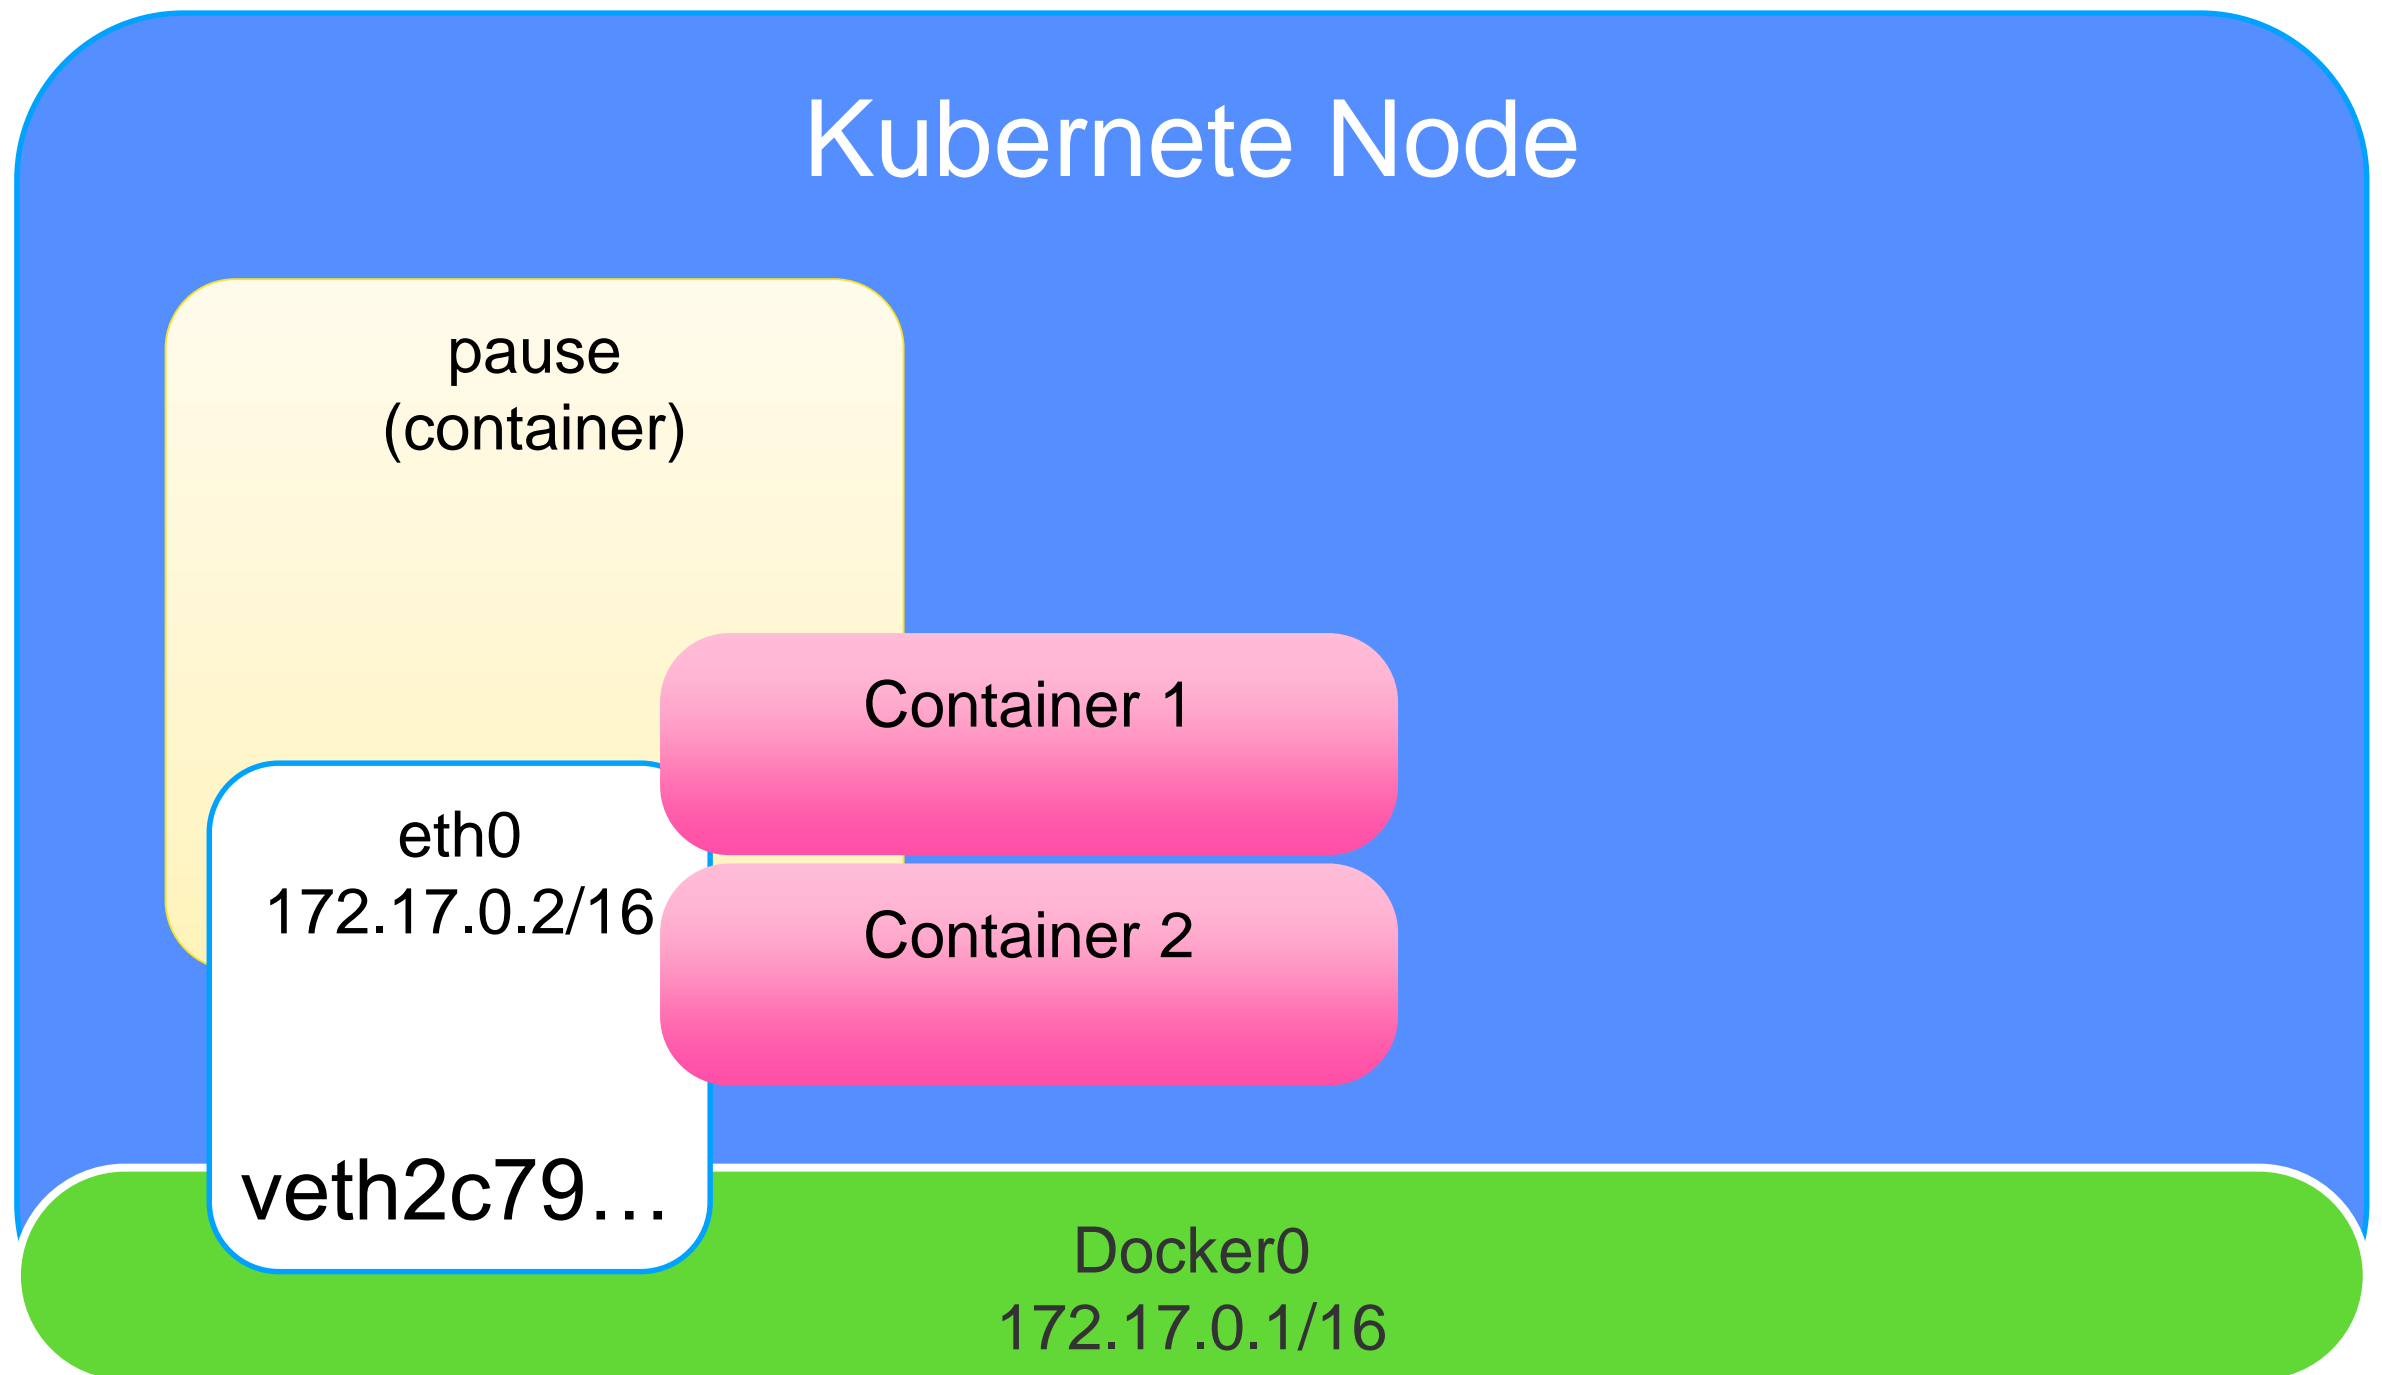
- Containers communication in a Pod
 - How to create a Pod
 - Assign a Pod a unique IP
- Pods on different nodes
- Pod to Service

Containers in a Pod



How to have many containers into a Pod?

Pause



Create pause, and “attach” containers to its network

Pod Networking

- Container unique IP -> Pod unique IP
 - The pause container get its **IP** and then pause
 - Add containers to pause's networks
 - Containers communicate with localhost
 - Containers share the same IP across cluster

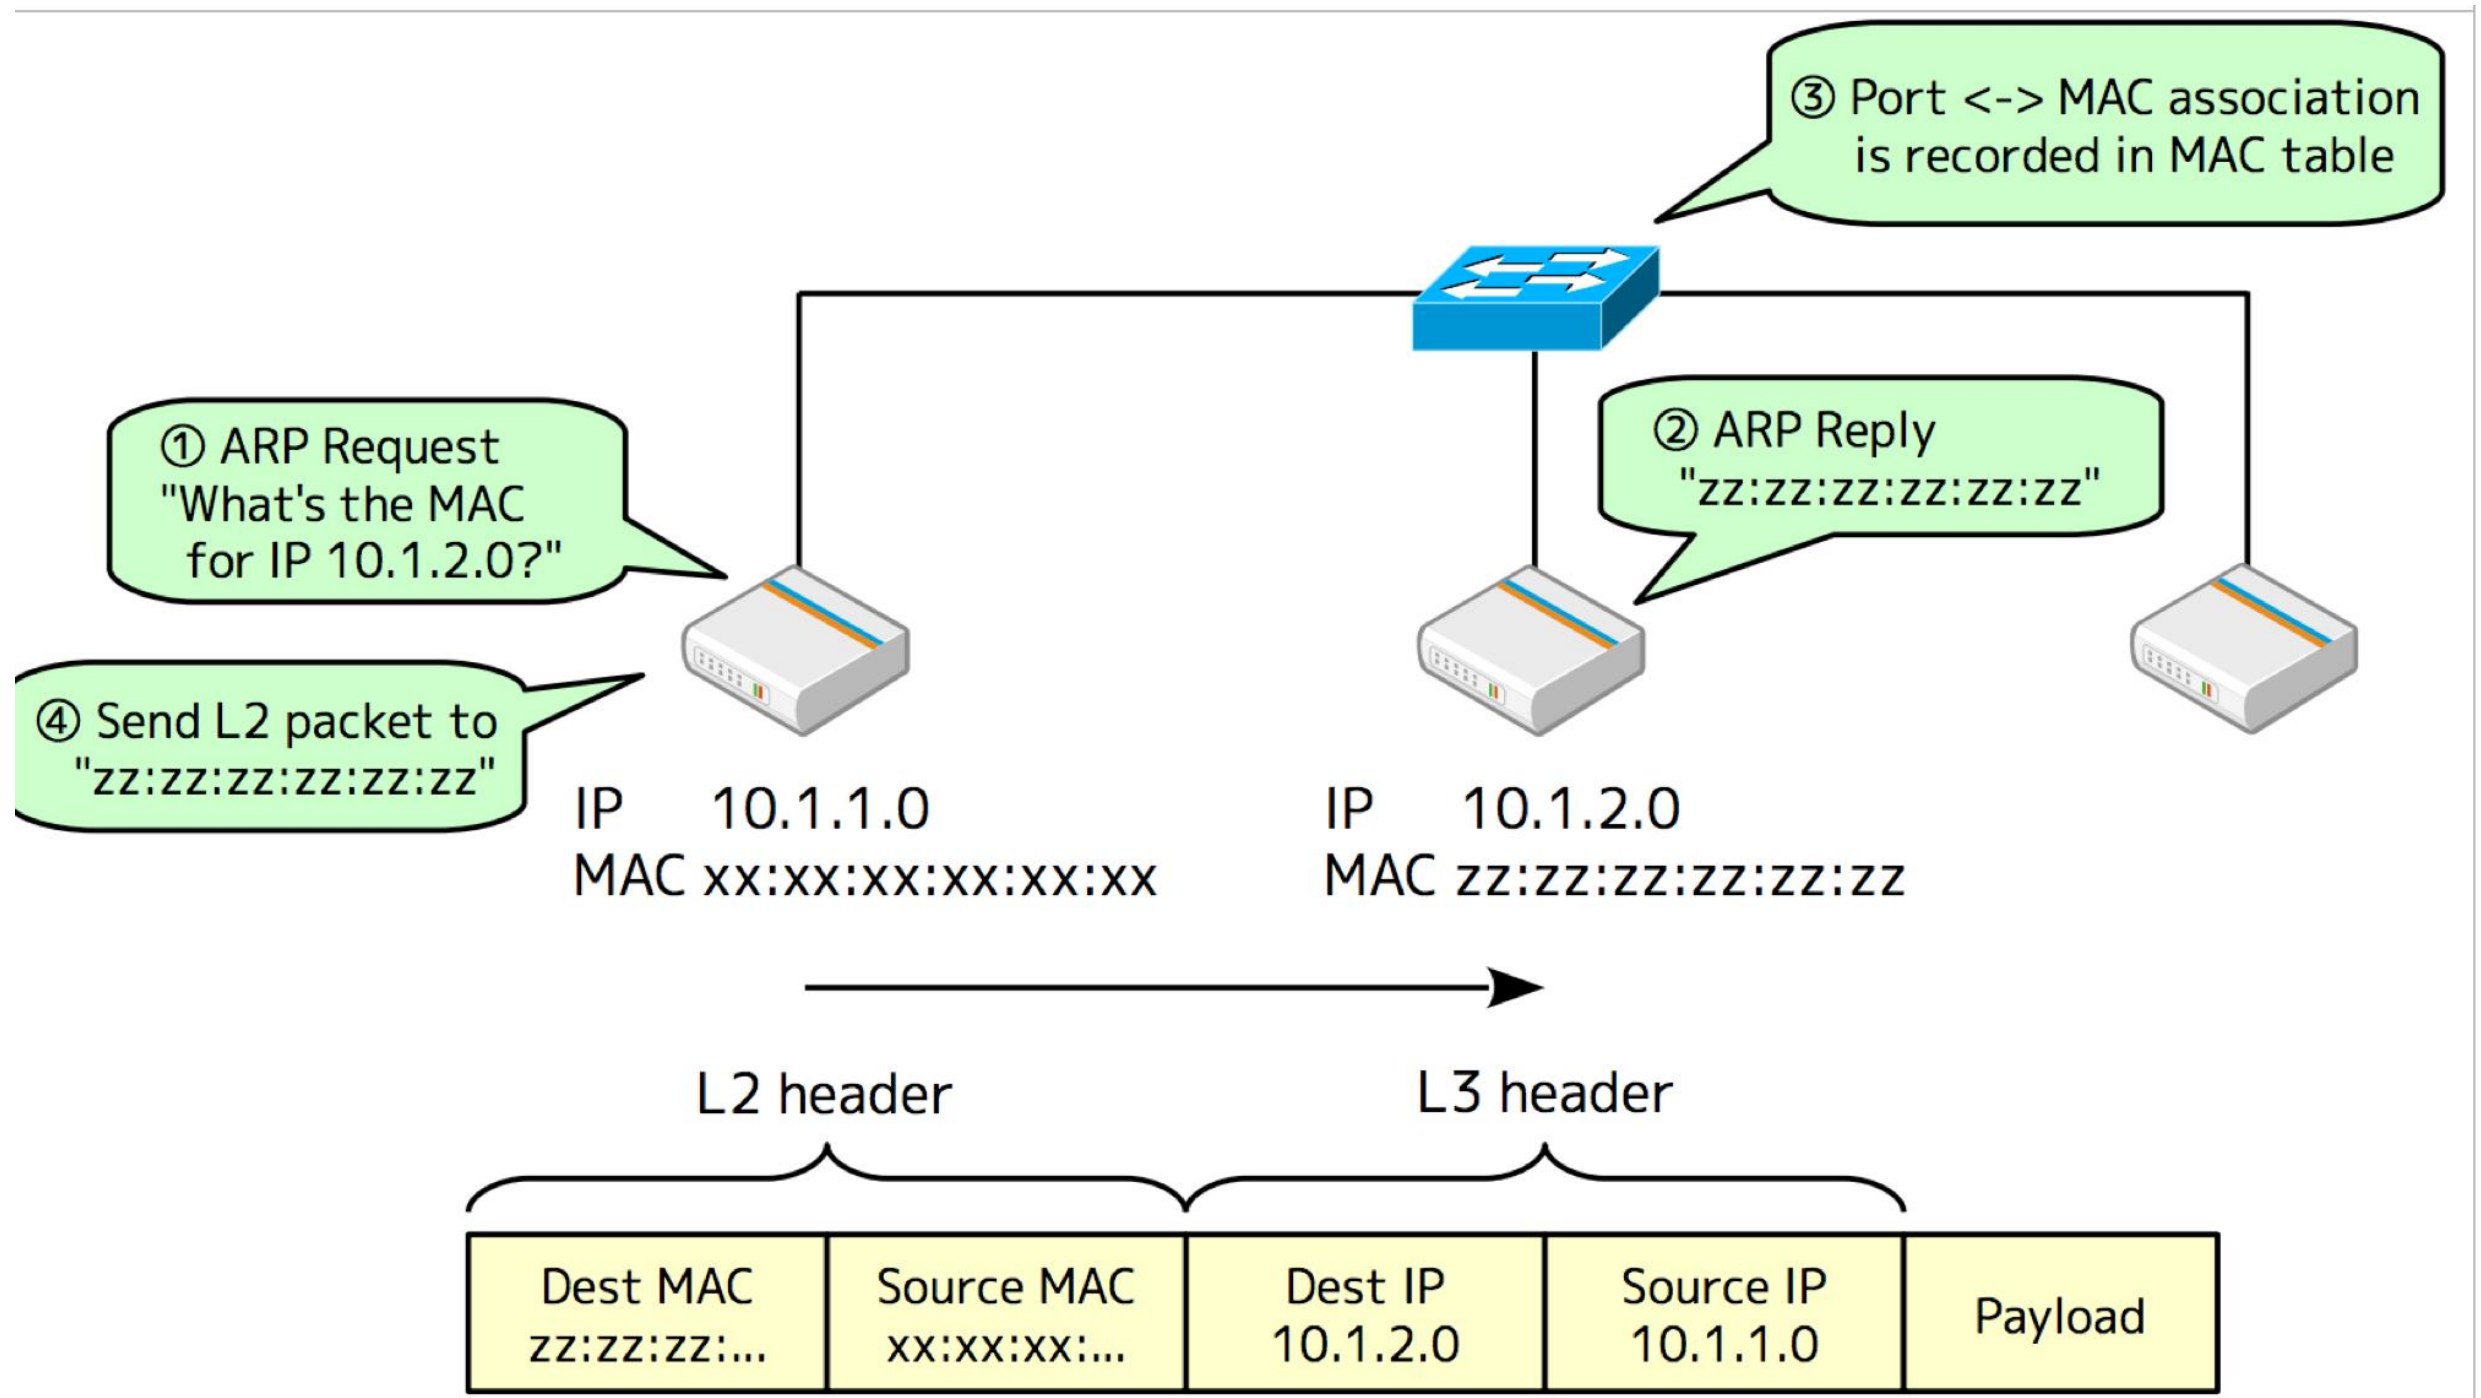
Kubernetes Networks

- Containers communication in a Pod
 - How to create a Pod
 - Assign a Pod a unique IP
- Pods on different nodes
- Pod to Service

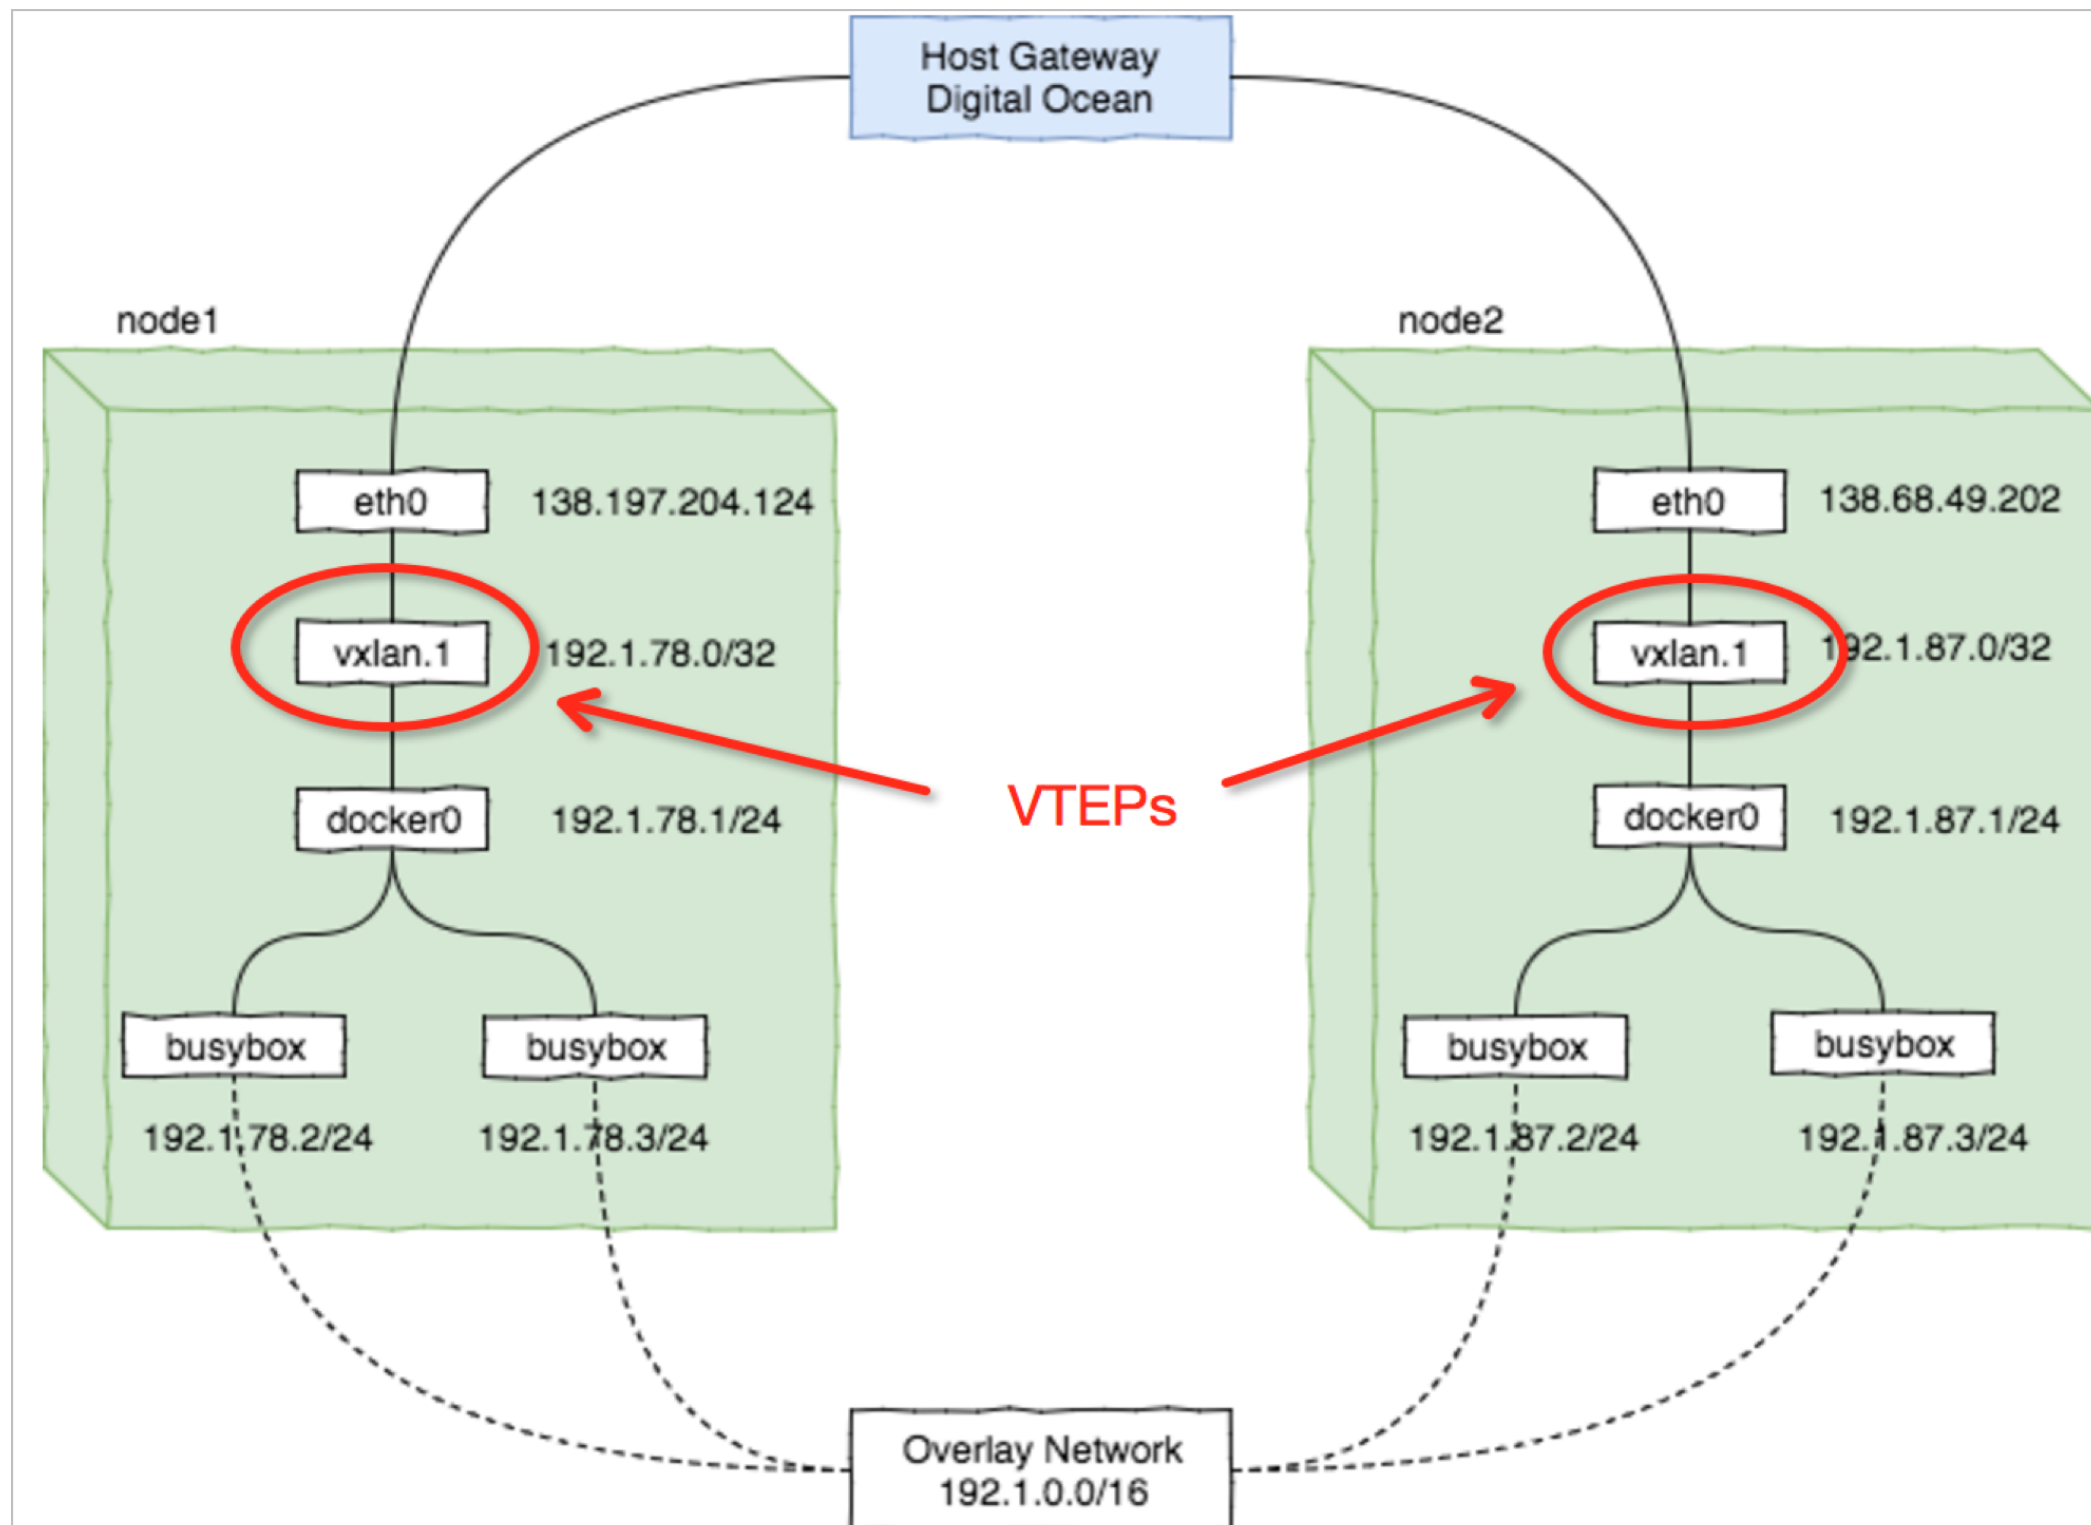
Flannel

- A flanneld on each host
- Flanneld creates a subnet for each host out of a larger address space
- Packet forward mechanism: VXLAN

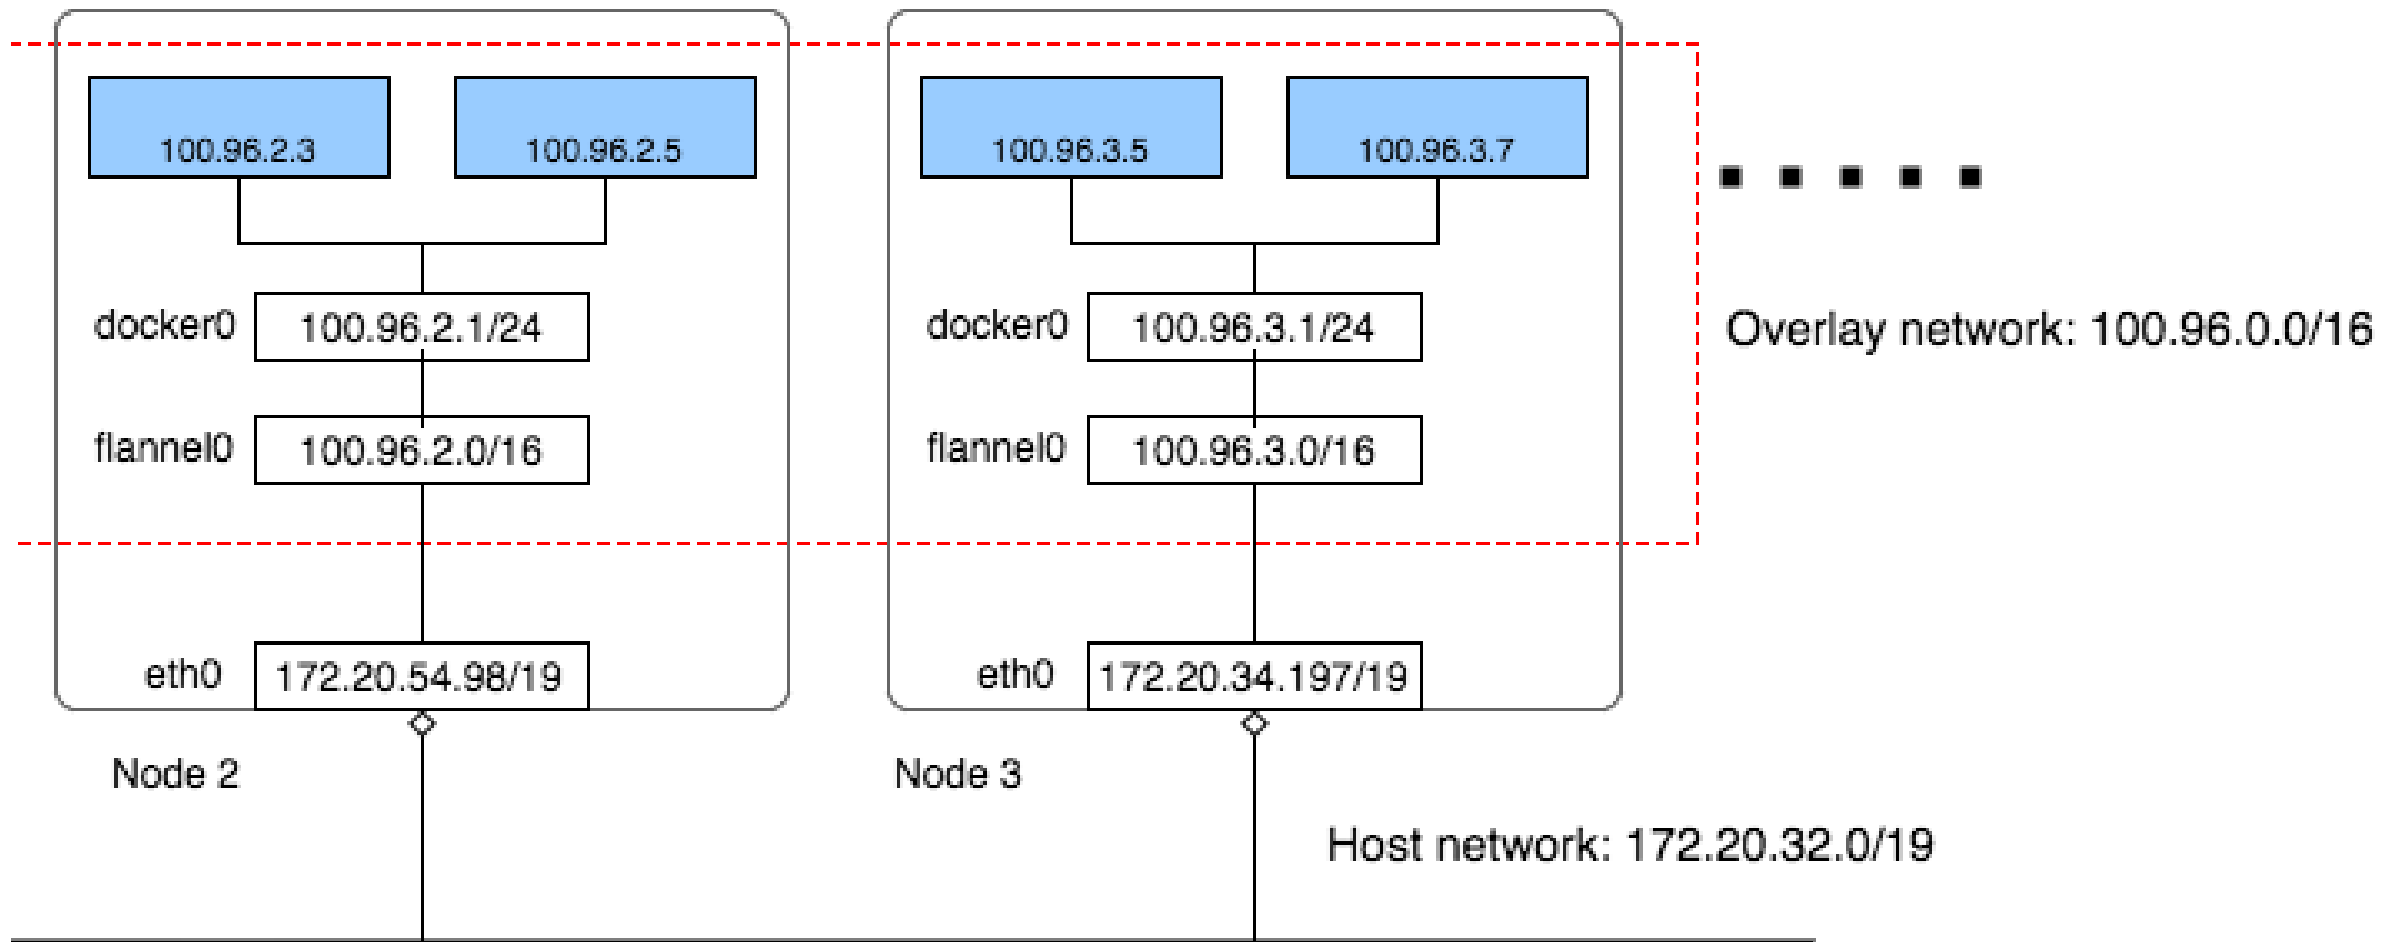
Virtual eXtensible Local Area Network



VXLAN



Overlay Network



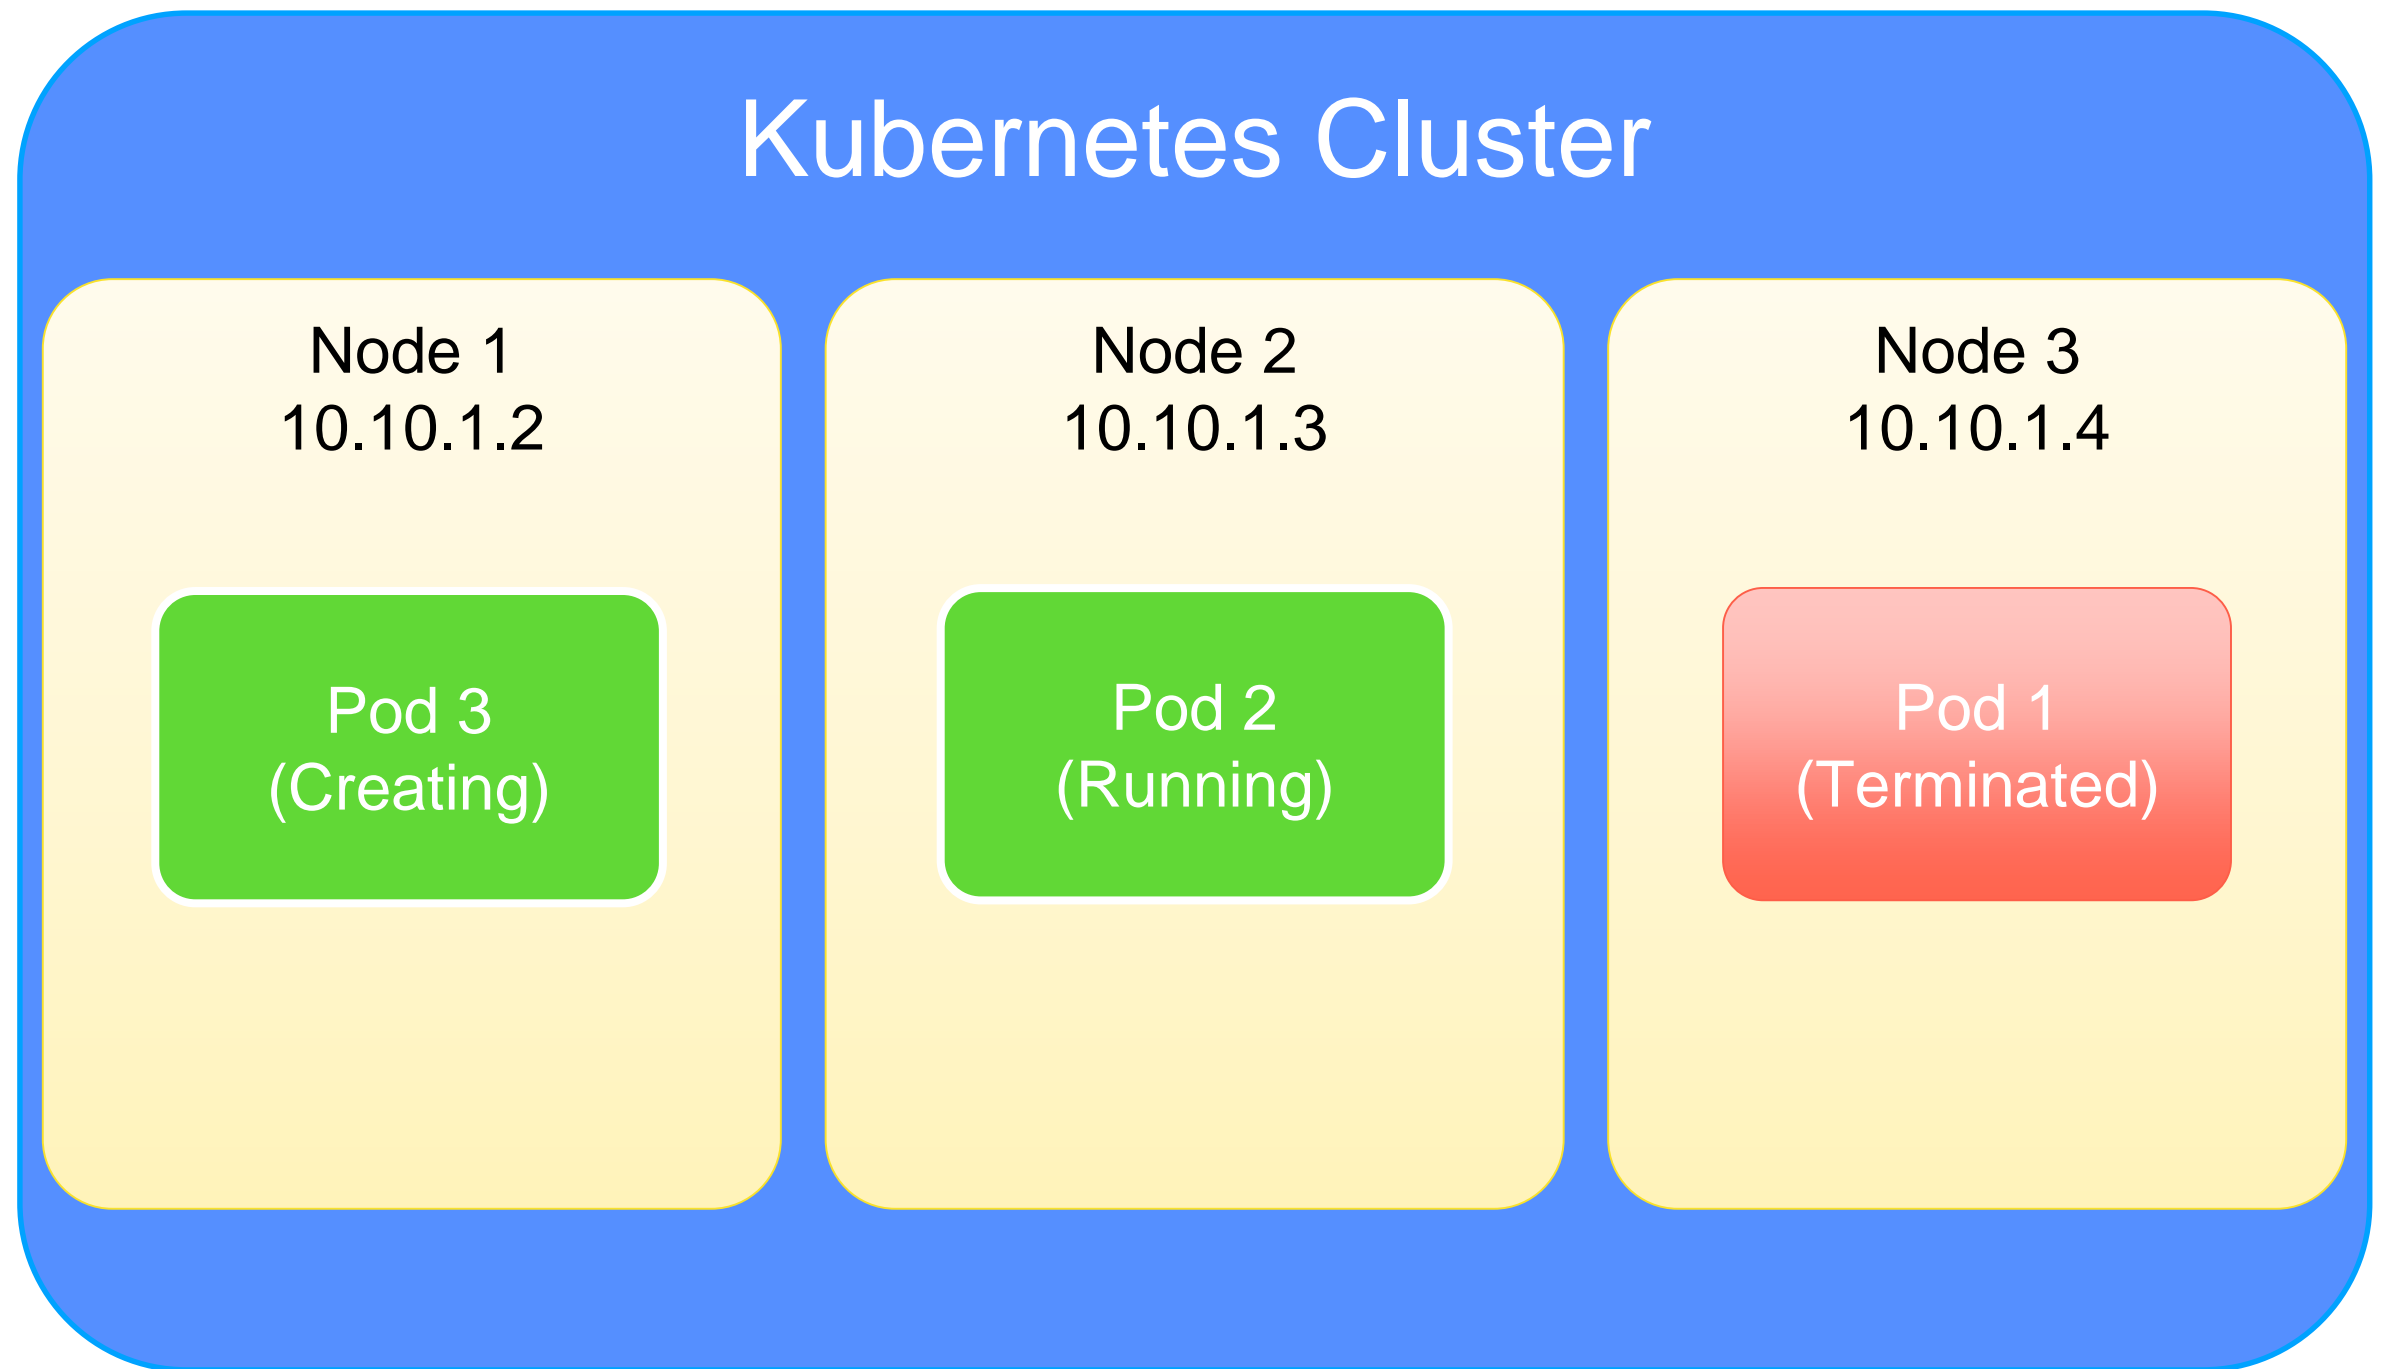
Kubernetes Networks

- Containers communication in a Pod -> localhost
 - How to create a Pod -> pause
 - Assign a Pod a unique IP -> flannel address space(CNI)
- Pods communicates across different nodes
-> flannel (vxlan, overlay networks)
- Pod to Service

Kubernetes Networks

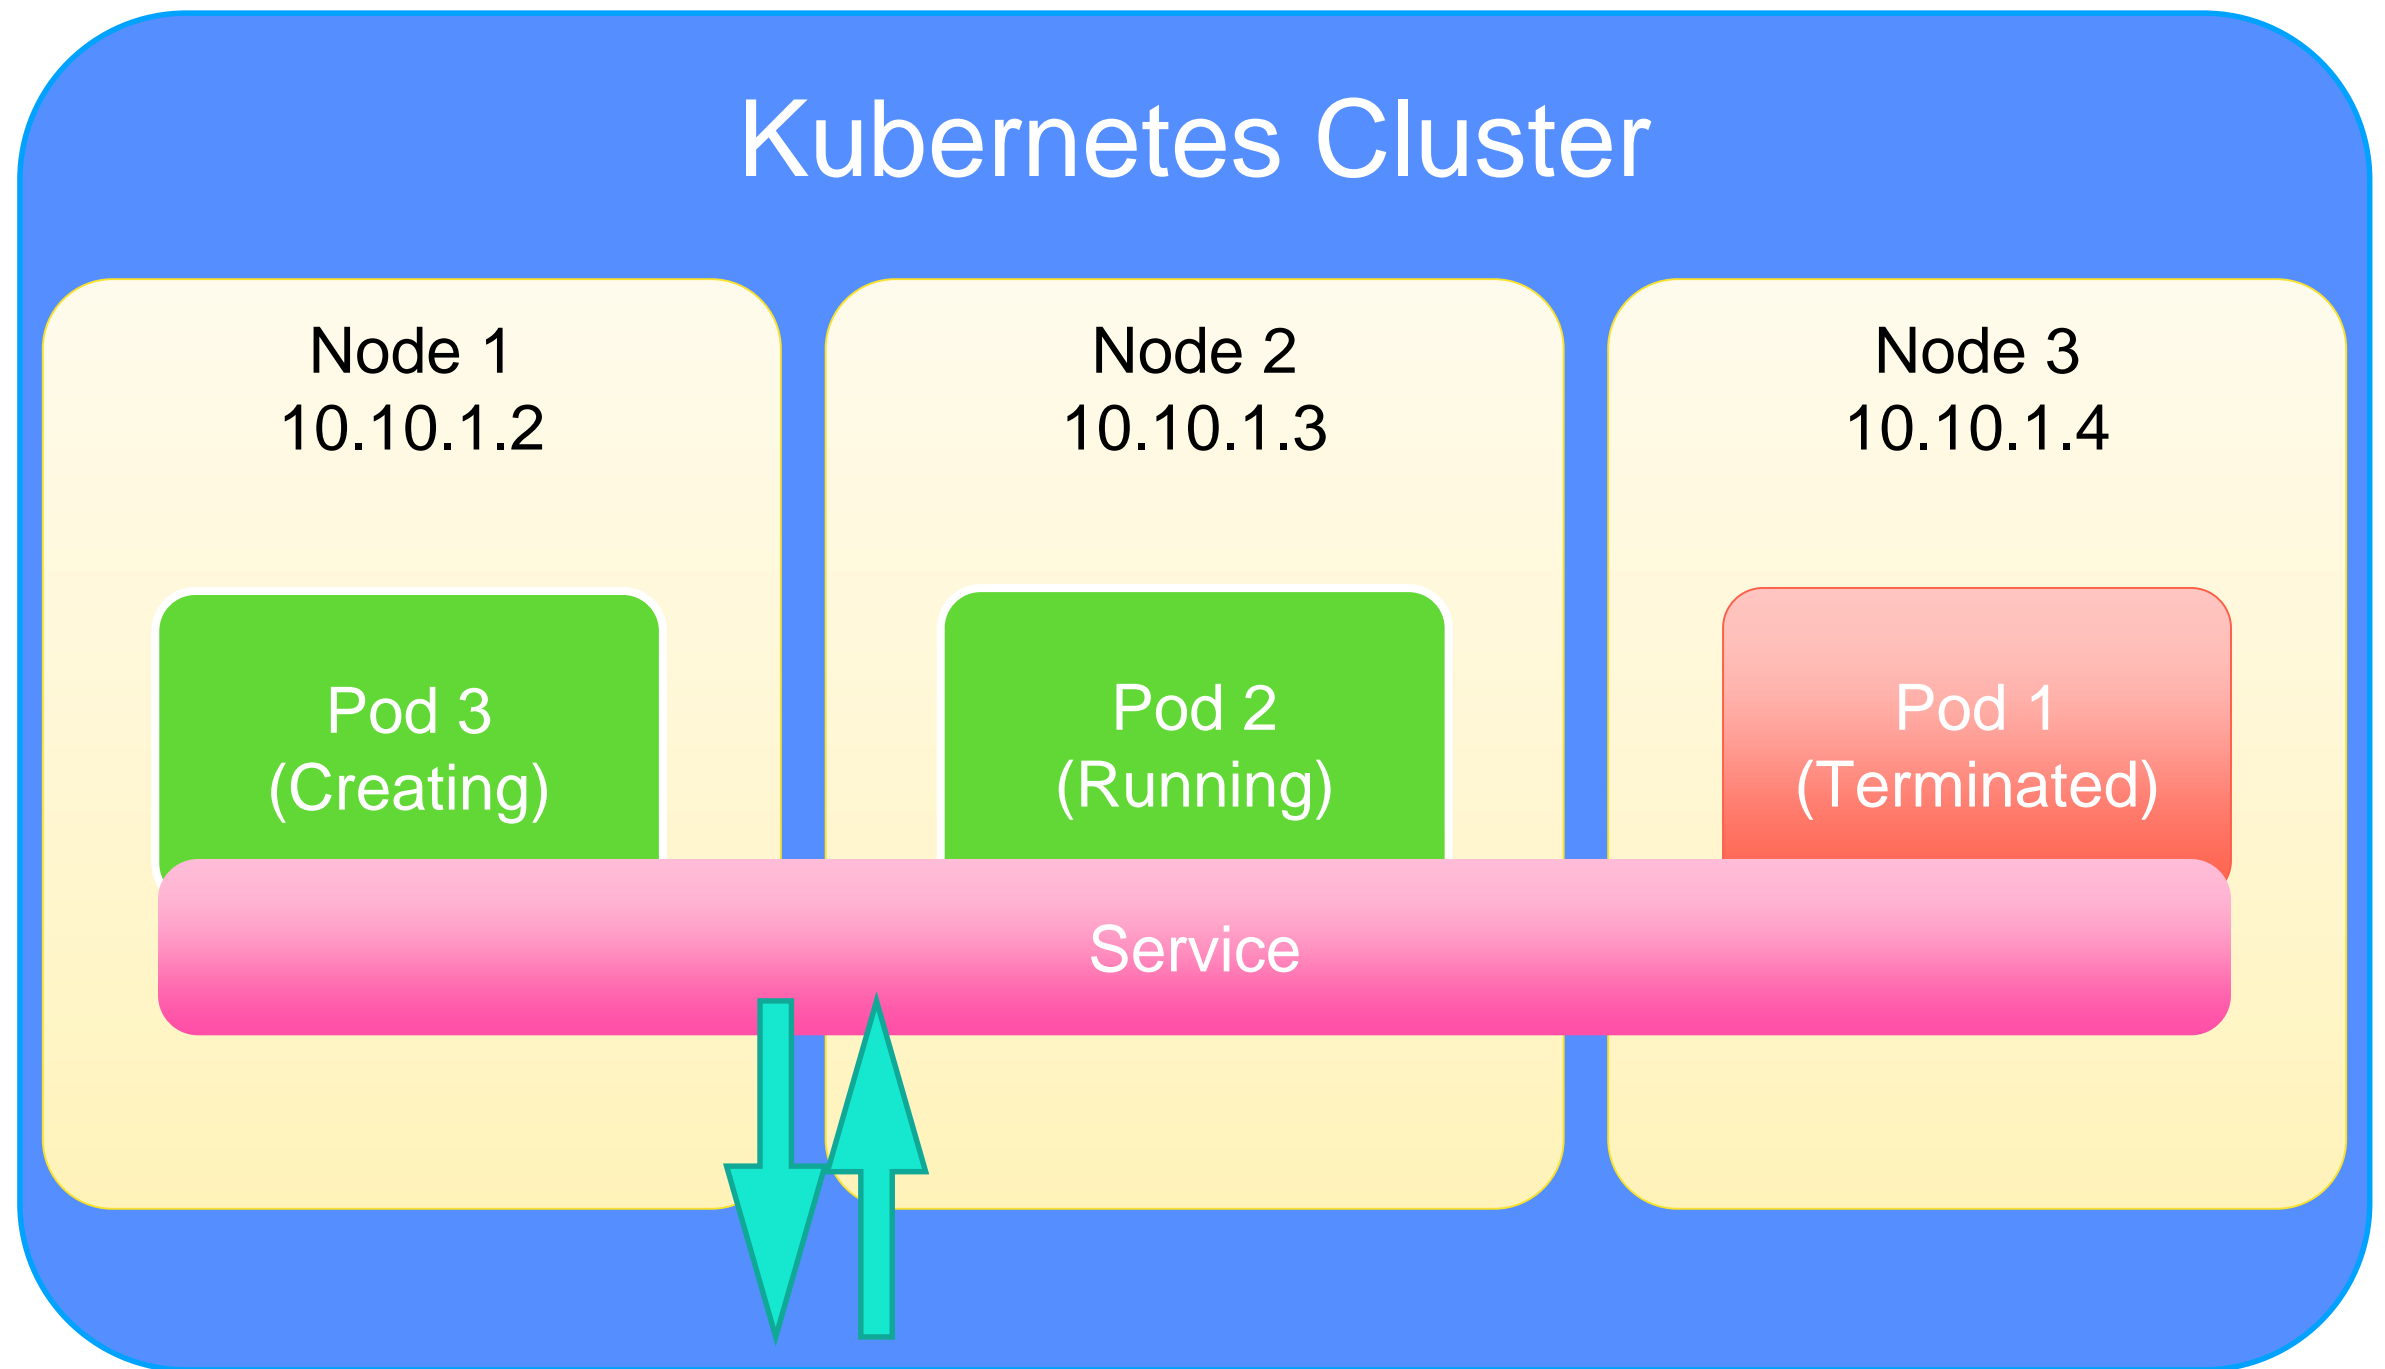
- Pod has unique IP but hard to access through pure IP
- Pod to kubernetes Service

Pod to Service



How to access app instead of access to a specific Pod IP?

Kubernetes Service

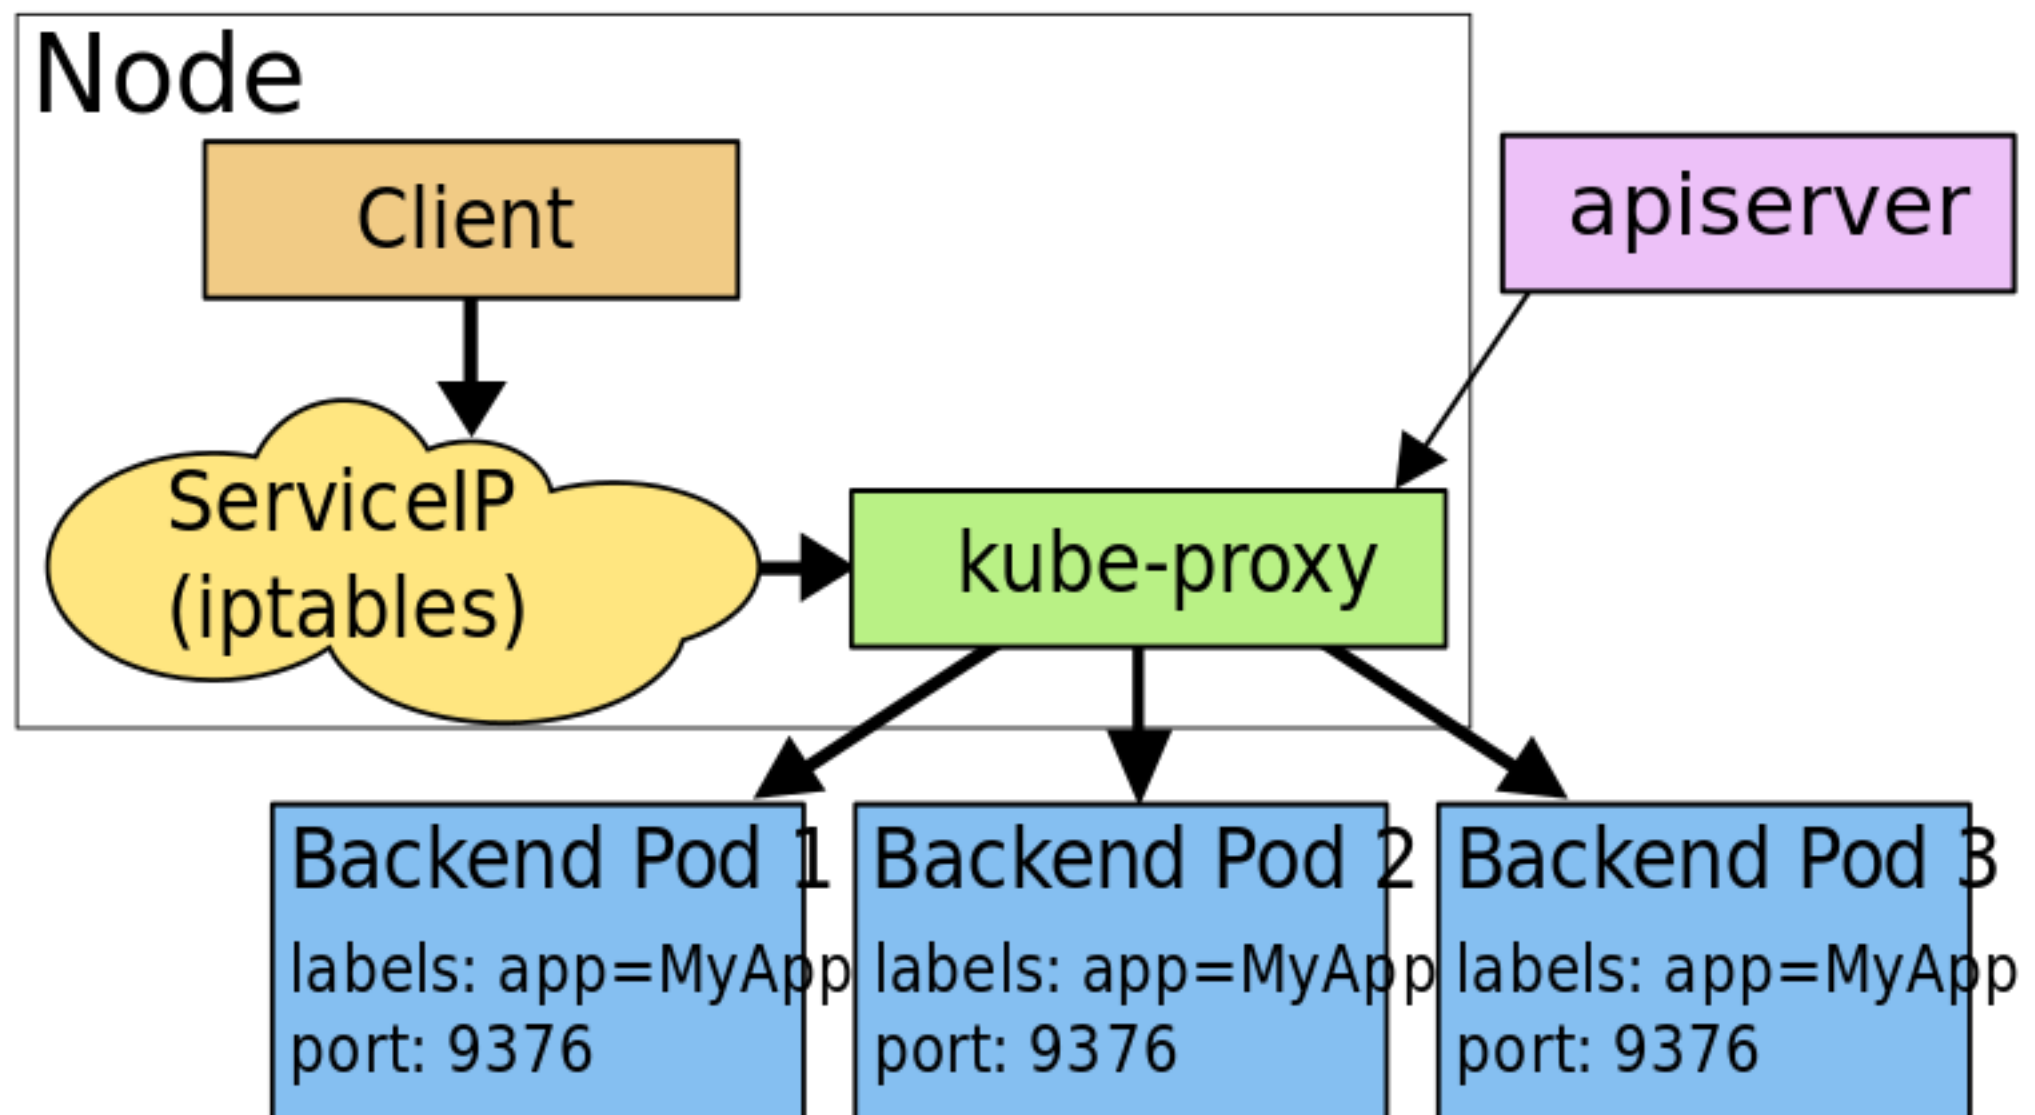


Access to a set of Pods through a single endpoint

Kubernetes Service

- A Kubernetes Service is an abstraction which defines a logical set of Pods and a policy by which to access them - sometimes called a micro-service.

Kube-proxy



Kube-proxy

- kube-proxy is responsible for implementing a form of virtual IP for Services of type other than ExternalName.
- In short, update iptable rules to nodes for each services

```
node$ sudo iptables-save -t nat
```

```
-A KUBE-SERVICES ! -s 10.233.64.0/18 -d 10.233.0.1/32 -p tcp -m comment --comment "default/kubernetes:https"
```

```
-A KUBE-SERVICES -d 10.233.0.1/32 -p tcp -m comment --comment "default/kubernetes:https"
```

Discovering Service

- Environment variables injection by kubelet
REDIS_MASTER_SERVICE_HOST=10.0.0.11
REDIS_MASTER_SERVICE_PORT=6379
- DNS (Recommended)
The DNS server watches the Kubernetes API for new Service and creates a set of DNS records for each

Kube-dns

- Kubernetes DNS schedules a DNS Pod and Service on the cluster, and configures the kubelets to tell individual containers to use the DNS Service's IP to resolve DNS names.

Service

`my-svc.my-namespace.svc.cluster.local -> nginx.default.svc.cluster.local`

Pod

`pod-ip-address.my-namespace.pod.cluster.local -> nginx-1.default.pod.cluster.local`

Kubernetes Networks

- Containers communication in a Pod
-> localhost
- Assign a Pod a unique IP
-> flannel address space
- Pods on different nodes
-> flannel (VXLAN, overlay networks)
- Pod to Service
-> service, proxy, dns

Thank you

Kubernetes Networks

Thank you